

```
In [1]: # rainfall prediction in australia for the next day using 10 years of data
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import StandardScaler
import tensorflow as tf
from tensorflow import keras
from keras.models import Sequential
from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
import seaborn as sns
import datetime
from sklearn.preprocessing import LabelEncoder
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import seaborn as sns
from keras.layers import Dense, BatchNormalization, Dropout, LSTM
from keras.models import Sequential
from keras.utils import to_categorical
from keras.optimizers import Adam
from tensorflow.keras import regularizers
from sklearn.metrics import precision_score, recall_score, confusion_matrix
from keras import callbacks
from sklearn.metrics import confusion_matrix
```

```
In [2]: data = pd.read_csv("/content/drive/MyDrive/weatherAUS.csv")
```

```
In [ ]: data.head(n=10)
```

```
Out[3]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpe
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	4
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	4
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	4
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	2
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	4
5	2008-12-06	Albury	14.6	29.7	0.2	NaN	NaN	WNW	5
6	2008-12-07	Albury	14.3	25.0	0.0	NaN	NaN	W	5
7	2008-12-08	Albury	7.7	26.7	0.0	NaN	NaN	W	3
8	2008-12-09	Albury	9.7	31.9	0.0	NaN	NaN	NNW	8
9	2008-12-10	Albury	13.1	30.1	1.4	NaN	NaN	W	2

10 rows × 23 columns

Basic Data cleaning and data analysis

```
In [3]: data.shape
```

```
Out[3]: (145460, 23)
```

```
In [42]:
```

```
In [5]: data.dtypes
```

```
Out[5]: Date                object
Location                   object
MinTemp                   float64
MaxTemp                   float64
Rainfall                  float64
Evaporation               float64
Sunshine                  float64
WindGustDir               object
WindGustSpeed             float64
WindDir9am                object
WindDir3pm                object
WindSpeed9am              float64
WindSpeed3pm              float64
Humidity9am               float64
Humidity3pm               float64
Pressure9am               float64
Pressure3pm               float64
Cloud9am                  float64
Cloud3pm                  float64
Temp9am                   float64
Temp3pm                   float64
RainToday                  object
RainTomorrow              object
dtype: object
```

```
In [6]: data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Date                  145460 non-null object  
 1   Location              145460 non-null object  
 2   MinTemp               143975 non-null float64 
 3   MaxTemp               144199 non-null float64 
 4   Rainfall              142199 non-null float64 
 5   Evaporation           82670 non-null  float64 
 6   Sunshine              75625 non-null  float64 
 7   WindGustDir           135134 non-null object  
 8   WindGustSpeed          135197 non-null float64 
 9   WindDir9am            134894 non-null object  
10   WindDir3pm            141232 non-null object  
11   WindSpeed9am          143693 non-null float64 
12   WindSpeed3pm          142398 non-null float64 
13   Humidity9am           142806 non-null float64 
14   Humidity3pm           140953 non-null float64 
15   Pressure9am           130395 non-null float64 
16   Pressure3pm           130432 non-null float64 
17   Cloud9am              89572 non-null  float64 
18   Cloud3pm              86102 non-null  float64 
19   Temp9am               143693 non-null float64 
20   Temp3pm               141851 non-null float64 
21   RainToday             142199 non-null object  
22   RainTomorrow          142193 non-null object  
dtypes: float64(16), object(7)
memory usage: 25.5+ MB
```

```
In [ ]:
```

we can see that we have null values in our dataset. Nevertheless first try to get some information from the dataset then will come back to the missing values.

```
In [ ]: # checking the class balance using simple graph.  
sns.set(style="white")  
sns.countplot(x= data["RainTomorrow"], palette=["lightgray", "green"] )  
sns.despine(top=True, right=True)  
plt.title("Rain Tomorrow Count")  
plt.xlabel("Rain Tomorrow")  
plt.ylabel("Count")
```

Out[9]: Text(0, 0.5, 'Count')



In []:

In []:

```
In [ ]: class_counts = data['RainTomorrow'].value_counts()

# Identify the majority and minority classes
majority_class = class_counts.idxmax()
minority_class = class_counts.idxmin()

# Calculate imbalance ratio
imbalance_ratio = class_counts[majority_class] / class_counts[minority_class]

print(f"Class Counts:\n{class_counts}")
print(f"Imbalance Ratio: {imbalance_ratio:.2f}")
```

```
Class Counts:
No      110316
Yes      31877
Name: RainTomorrow, dtype: int64
Imbalance Ratio: 3.46
```

```
In [ ]:
```

Such unbalanced class ratios can have an effect on machine learning methods and model performance.

strategies to deal with unbalanced classes:

Resampling: To balance the class distribution, you might oversample the minority class or undersample the dominant class.

Synthetic Data Generation: Methods such as SMOTE can be used to create synthetic instances of the minority class.

Cost-Sensitive Learning: Change the algorithm's learning process to penalise minority misclassification more severely.

Algorithm Selection: Select algorithms that are built to deal with skewed data, such as ensemble methods or algorithms with built-in class weights.

```
In [ ]:
```

We will use ANN as our machine learning method because it works well with unbalanced data as it contains built-in class weights to deal with unbalanced data.

```
In [ ]:
```

Covertng date to date-time format

```
In [7]: data['Date'] = pd.to_datetime(data["Date"])
data['year'] = data.Date.dt.year
data['month'] = data.Date.dt.month
data['day'] = data.Date.dt.day
```

```
In [8]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 26 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Date                  145460 non-null  datetime64[ns]
 1   Location              145460 non-null  object
 2   MinTemp               143975 non-null  float64
 3   MaxTemp               144199 non-null  float64
 4   Rainfall              142199 non-null  float64
 5   Evaporation           82670 non-null   float64
 6   Sunshine              75625 non-null   float64
 7   WindGustDir           135134 non-null  object
 8   WindGustSpeed          135197 non-null  float64
 9   WindDir9am            134894 non-null  object
10   WindDir3pm            141232 non-null  object
11   WindSpeed9am          143693 non-null  float64
12   WindSpeed3pm          142398 non-null  float64
13   Humidity9am           142806 non-null  float64
14   Humidity3pm           140953 non-null  float64
15   Pressure9am           130395 non-null  float64
16   Pressure3pm           130432 non-null  float64
17   Cloud9am              89572 non-null   float64
18   Cloud3pm              86102 non-null   float64
19   Temp9am               143693 non-null  float64
20   Temp3pm               141851 non-null  float64
21   RainToday             142199 non-null  object
22   RainTomorrow          142193 non-null  object
23   year                  145460 non-null  int64
24   month                 145460 non-null  int64
25   day                   145460 non-null  int64
dtypes: datetime64[ns](1), float64(16), int64(3), object(6)
memory usage: 28.9+ MB
```

```
In [ ]: data.head()
```

```
Out[13]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpe
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	4
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	4
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	4
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	2
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	4

5 rows × 26 columns

```
In [9]:
```

```
Out[9]: year
2007    25.086885
2008    22.874359
2009    23.251019
2010    22.571247
2011    22.540180
2012    22.311424
2013    23.290936
2014    23.792739
2015    23.515652
2016    23.419984
2017    25.031702
Name: MaxTemp, dtype: float64
```

```
In [9]:
```

```
In [ ]: yearly_avg_rain = data.groupby('year')['Rainfall'].mean()
yearly_avg_rain
```

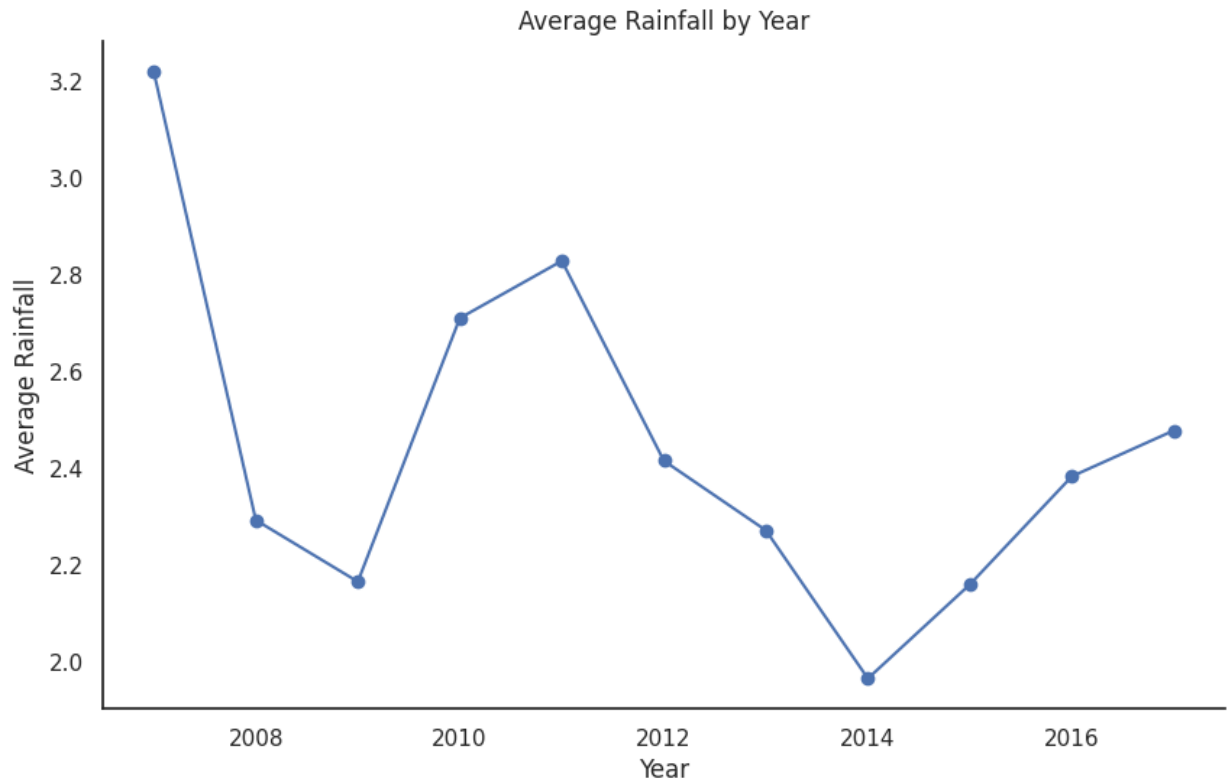
```
Out[16]: year
2007    3.219672
2008    2.293541
2009    2.166385
2010    2.710924
2011    2.829197
2012    2.416200
2013    2.272402
2014    1.966341
2015    2.160753
2016    2.384054
2017    2.478834
Name: Rainfall, dtype: float64
```



```
In [ ]: yearly_avg_rain_df = yearly_avg_rain.reset_index()

plt.figure(figsize=(10, 6))
plt.plot(yearly_avg_rain_df['year'], yearly_avg_rain_df['Rainfall'], marker
sns.despine(top=True, right=True)
plt.title('Average Rainfall by Year')
plt.xlabel('Year')
plt.ylabel('Average Rainfall')

plt.show()
```



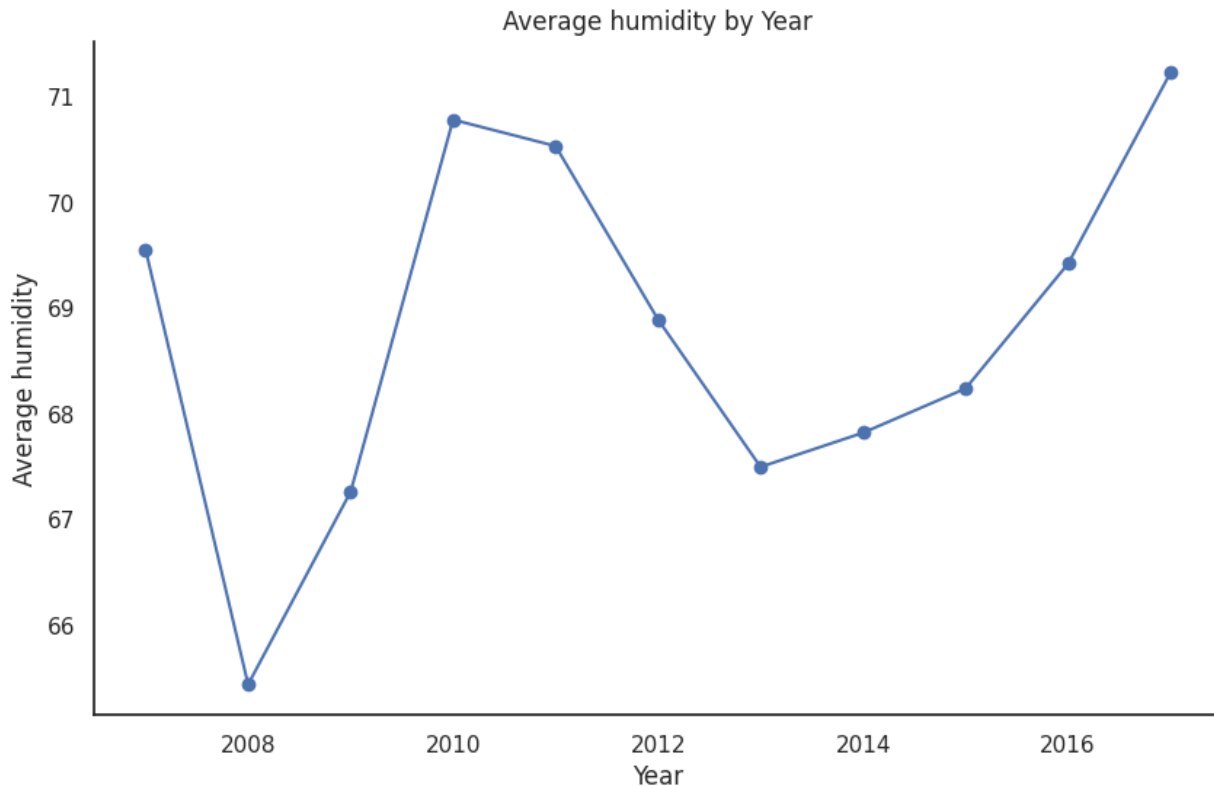
In [9]:

In [9]:

```
In [ ]: yearly_avg_humidity = data.groupby('year')['Humidity9am'].mean()  
yearly_avg_humidity
```

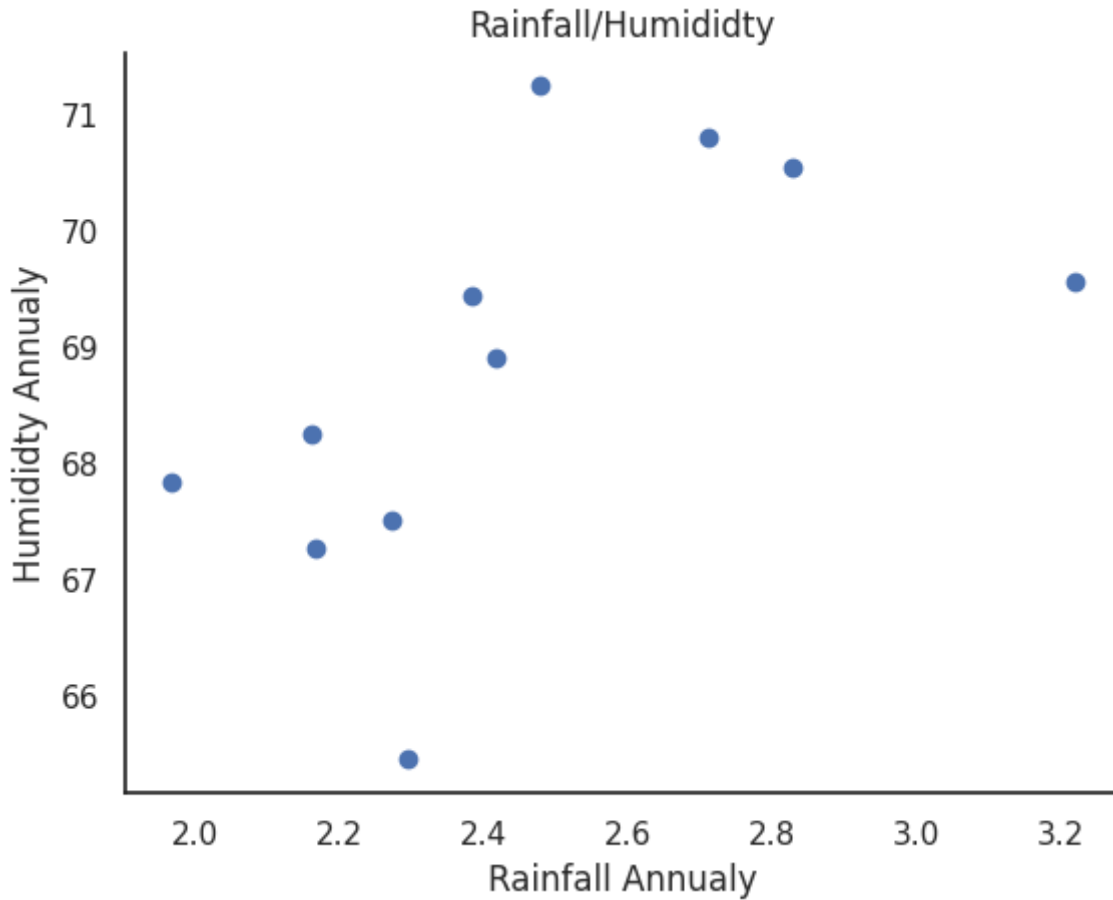
```
Out[20]: year  
2007    69.557377  
2008    65.442762  
2009    67.266126  
2010    70.788379  
2011    70.539129  
2012    68.896674  
2013    67.501632  
2014    67.825234  
2015    68.244000  
2016    69.424885  
2017    71.234636  
Name: Humidity9am, dtype: float64
```

```
In [ ]: yearly_avg_humidity_df = yearly_avg_humidity.reset_index()  
  
plt.figure(figsize=(10, 6))  
plt.plot(yearly_avg_humidity_df['year'], yearly_avg_humidity_df['Humidity9a  
sns.despine(top=True, right=True)  
plt.title('Average humidity by Year')  
plt.xlabel('Year')  
plt.ylabel('Average humidity')  
  
plt.show()
```



```
In [ ]: plt.scatter(yearly_avg_rain, yearly_avg_humidity )
sns.despine(top=True, right=True)
plt.title('Rainfall/Humidity')
plt.xlabel('Rainfall Annualy')
plt.ylabel('Humididty Annualy')
```

```
Out[22]: Text(0, 0.5, 'Humididty Annualy')
```



```
In [11]:
```

Type *Markdown* and LaTeX: α^2

The month and year columns must be converted to cyclic format. This allows the neural network to capture cyclical patterns in the data without introducing discontinuities.

```
In [12]: # Convert month and year to cyclic representations
def encode_cyclic(value, period):
    return np.sin(2 * np.pi * value / period), np.cos(2 * np.pi * value / p

data[['year_sin', 'year_cos']] = data['year'].apply(lambda x: pd.Series(enc
data[['month_sin', 'month_cos']] = data['month'].apply(lambda x: pd.Series(
data[['day_sin', 'day_cos']] = data['day'].apply(lambda x: pd.Series(encode
```

```
In [ ]:
```

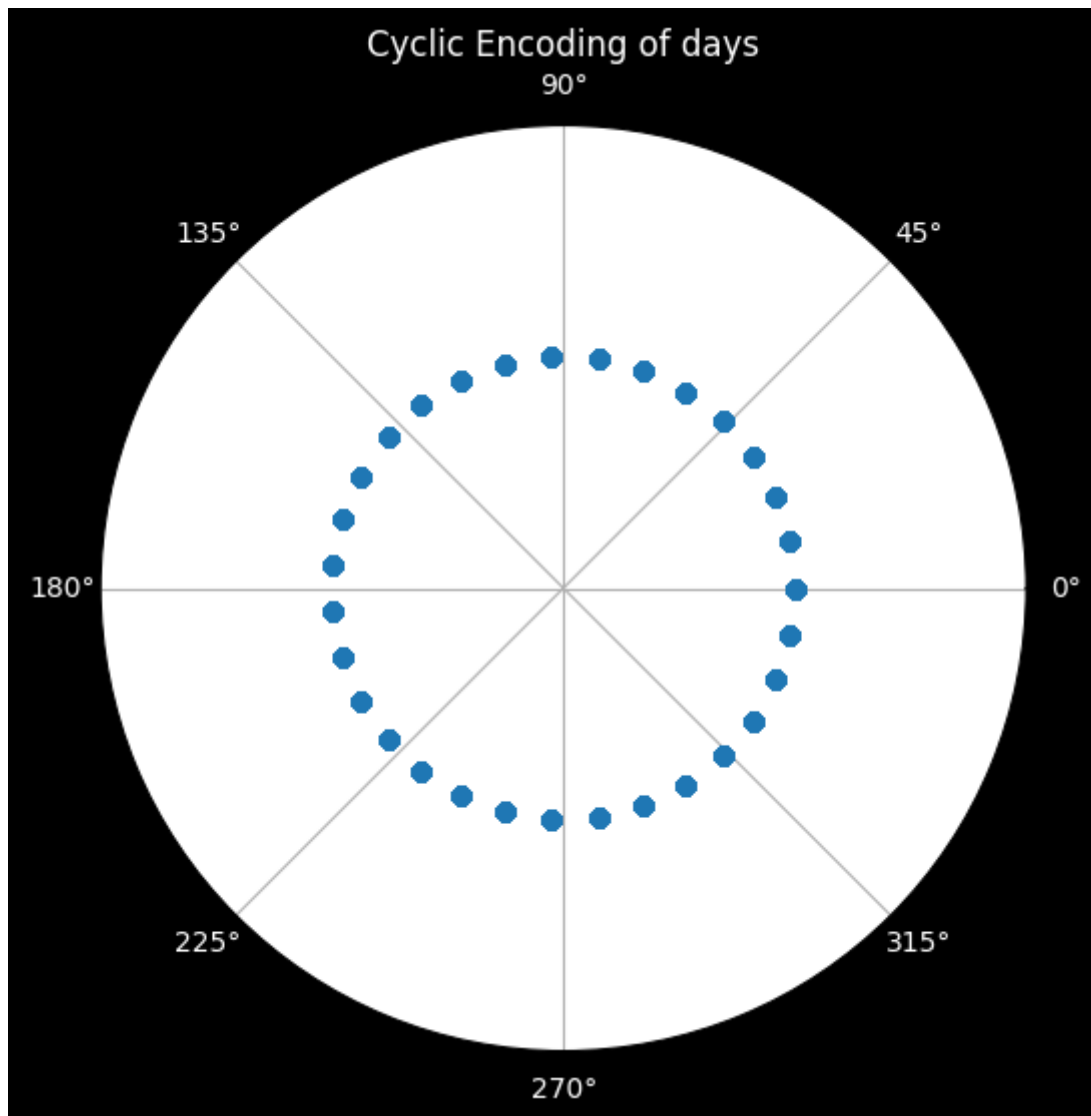
```
In [ ]:
```

```
In [ ]:
```

```
In [13]: fig = plt.figure(figsize=(6, 6))
fig.patch.set_facecolor('black')
ax = fig.add_subplot(111, projection='polar')
ax.plot(np.arctan2(data['day_sin'], data['day_cos']), np.sqrt(np.square(dat


ax.set_title("Cyclic Encoding of days", color='white') # Set plot title co
ax.tick_params(axis='both', colors='white')
ax.set_rticks([]) # Hide radial tick labels

plt.show()
```




```
In [14]: fig = plt.figure(figsize=(6, 6))
fig.patch.set_facecolor('black')
ax = fig.add_subplot(111, projection='polar')
ax.plot(np.arctan2(data['month_sin'], data['month_cos']), np.sqrt(np.square
ax.set_title("Cyclic Encoding of Month", color='white') # Set plot title c
ax.tick_params(axis='both', colors='white')
ax.set_rticks([]) # Hide radial tick labels

plt.show()
```



```
In [15]: fig = plt.figure(figsize=(6, 6))
fig.patch.set_facecolor('black')
ax = fig.add_subplot(111, projection='polar')
ax.plot(np.arctan2(data['year_sin'], data['year_cos']), np.sqrt(np.square(d
ax.set_title("Cyclic Encoding of Month", color='white') # Set plot title c
ax.tick_params(axis='both', colors='white')
ax.set_rticks([]) # Hide radial tick labels

plt.show()
```



```
In [16]: data.isnull().sum()
```

```
Out[16]: Date                0
Location                    0
MinTemp                    1485
MaxTemp                    1261
Rainfall                   3261
Evaporation                62790
Sunshine                   69835
WindGustDir                10326
WindGustSpeed              10263
WindDir9am                 10566
WindDir3pm                 4228
WindSpeed9am               1767
WindSpeed3pm               3062
Humidity9am                2654
Humidity3pm                4507
Pressure9am                15065
Pressure3pm                15028
Cloud9am                   55888
Cloud3pm                   59358
Temp9am                    1767
Temp3pm                    3609
RainToday                  3261
RainTomorrow               3267
year                       0
month                     0
day                       0
year_sin                   0
year_cos                   0
month_sin                  0
month_cos                  0
day_sin                    0
day_cos                    0
dtype: int64
```

```
In [ ]:
```

As we see missing values in our dataset, we will replace them with mode for categorical variables and meadian for numeric values because mean is heavily influenced by outliers.

```
In [17]: categorical_missing = (data.dtypes == "object")
object_cols = list(categorical_missing[categorical_missing].index)
```

```
In [18]: for i in object_cols:
          data[i].fillna(data[i].mode()[0], inplace=True)
```



```
In [19]: for i in object_cols:
          print(i, data[i].isnull().sum())
```

```
Location 0
WindGustDir 0
WindDir9am 0
WindDir3pm 0
RainToday 0
RainTomorrow 0
```

```
In [20]: numerical_missing = (data.dtypes == "float64")
          num_cols = list(numerical_missing[numerical_missing].index)
```


```
In [ ]:
```

```
In [21]: for i in num_cols:
          data[i].fillna(data[i].median(), inplace=True)

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 32 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  145460 non-null  datetime64[ns]
1   Location              145460 non-null  object
2   MinTemp               145460 non-null  float64
3   MaxTemp               145460 non-null  float64
4   Rainfall              145460 non-null  float64
5   Evaporation           145460 non-null  float64
6   Sunshine              145460 non-null  float64
7   WindGustDir           145460 non-null  object
8   WindGustSpeed         145460 non-null  float64
9   WindDir9am            145460 non-null  object
10  WindDir3pm            145460 non-null  object
11  WindSpeed9am          145460 non-null  float64
12  WindSpeed3pm          145460 non-null  float64
13  Humidity9am           145460 non-null  float64
14  Humidity3pm           145460 non-null  float64
15  Pressure9am           145460 non-null  float64
16  Pressure3pm           145460 non-null  float64
17  Cloud9am              145460 non-null  float64
18  Cloud3pm              145460 non-null  float64
19  Temp9am               145460 non-null  float64
20  Temp3pm               145460 non-null  float64
21  RainToday             145460 non-null  object
22  RainTomorrow          145460 non-null  object
23  year                  145460 non-null  int64
24  month                 145460 non-null  int64
25  day                   145460 non-null  int64
26  year_sin              145460 non-null  float64
27  year_cos              145460 non-null  float64
28  month_sin             145460 non-null  float64
29  month_cos             145460 non-null  float64
30  day_sin               145460 non-null  float64
31  day_cos               145460 non-null  float64
dtypes: datetime64[ns](1), float64(22), int64(3), object(6)
memory usage: 35.5+ MB
```

```
In [22]: sns.boxenplot(data = data)
plt.xticks(rotation=90)
plt.show()
```



```
In [ ]:
```

As we can see, our variables have a wide range, making it difficult to spot outliers before scaling the data. As we have string values (city names), we will utilise label encoding.

```
In [ ]:
```

In [23]:

```
label_encoder = LabelEncoder()
for i in object_cols:
    data[i] = label_encoder.fit_transform(data[i])

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 32 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  145460 non-null  datetime64[ns]
1   Location              145460 non-null  int64
2   MinTemp               145460 non-null  float64
3   MaxTemp               145460 non-null  float64
4   Rainfall              145460 non-null  float64
5   Evaporation           145460 non-null  float64
6   Sunshine              145460 non-null  float64
7   WindGustDir           145460 non-null  int64
8   WindGustSpeed         145460 non-null  float64
9   WindDir9am            145460 non-null  int64
10  WindDir3pm            145460 non-null  int64
11  WindSpeed9am          145460 non-null  float64
12  WindSpeed3pm          145460 non-null  float64
13  Humidity9am           145460 non-null  float64
14  Humidity3pm           145460 non-null  float64
15  Pressure9am           145460 non-null  float64
16  Pressure3pm           145460 non-null  float64
17  Cloud9am              145460 non-null  float64
18  Cloud3pm              145460 non-null  float64
19  Temp9am               145460 non-null  float64
20  Temp3pm               145460 non-null  float64
21  RainToday             145460 non-null  int64
22  RainTomorrow          145460 non-null  int64
23  year                  145460 non-null  int64
24  month                 145460 non-null  int64
25  day                   145460 non-null  int64
26  year_sin              145460 non-null  float64
27  year_cos              145460 non-null  float64
28  month_sin             145460 non-null  float64
29  month_cos             145460 non-null  float64
30  day_sin               145460 non-null  float64
31  day_cos               145460 non-null  float64
dtypes: datetime64[ns](1), float64(22), int64(9)
memory usage: 35.5 MB
```


In [23]:

```
In [24]: features = data.drop(['RainTomorrow', 'day', 'month', 'year', 'Date'], axis
target = data['RainTomorrow']
```

```
In [25]: col_names = list(features.columns)
robust = RobustScaler()
features = robust.fit_transform(features)
features = pd.DataFrame(features, columns=col_names)
```

Weather variables can have outliers due to extreme weather events in the context of weather data prediction. In such circumstances, robust scaling may be a better option to avoid outliers skewing the training process.

```
In [27]: sns.boxenplot(data = features)
plt.xticks(rotation=90)
plt.show()
```



In [28]:

```

# Function to remove outliers using IQR method
def remove_outliers_iqr(features, threshold=1.5):
    for col in features.columns:
        Q1 = features[col].quantile(0.25)
        Q3 = features[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - threshold * IQR
        upper_bound = Q3 + threshold * IQR
        features = features[(features[col] >= lower_bound) & (features[col]
        return features


# Remove outliers using IQR method
df_no_outliers = remove_outliers_iqr(features)

```

```

In [29]: sns.boxenplot(data = df_no_outliers)
plt.xticks(rotation=90)
plt.show()

```



```

In [ ]: len(X_train.columns)

```

```

Out[103]: 27

```

In [31]:

In [34]:

```
X_train, X_test, y_train, y_test = train_test_split(features, target, test_
```

Early stopping is a technique used to prevent overfitting by monitoring the validation loss during training. If the validation loss starts increasing or stagnates after a certain number of epochs, early stopping halts the training process. This prevents the model from continuing to learn on the training data past the point where it starts to overfit.

```
In [35]: early_stopping = callbacks.EarlyStopping(
    min_delta=0.001,
    patience=20,
    restore_best_weights=True,
)
```


```
In [36]: model2 = Sequential()
model2.add(Dense(units = 27 , kernel_initializer = 'uniform', activation = 'relu'))
model2.add(Dense(units = 27, kernel_initializer = 'uniform', activation = 'relu'))
model2.add(Dense(units = 13, kernel_initializer='uniform', activation = 'relu'))
model2.add(Dense(units = 7, kernel_initializer = 'uniform', activation = 'relu'))
model2.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))
opt = Adam(learning_rate=0.00009)
model2.compile(optimizer = opt, loss = 'binary_crossentropy', metrics = ['accuracy'])
history2=model2.fit(X_train, y_train, epochs = 50 ,callbacks=[early_stopping])
```

```
Epoch 1/50
2910/2910 [=====] - 10s 3ms/step - loss: 0.4553 - accuracy: 0.8113 - val_loss: 0.3841 - val_accuracy: 0.8371
Epoch 2/50
2910/2910 [=====] - 6s 2ms/step - loss: 0.3719 - accuracy: 0.8423 - val_loss: 0.3683 - val_accuracy: 0.8426
Epoch 3/50
2910/2910 [=====] - 9s 3ms/step - loss: 0.3620 - accuracy: 0.8453 - val_loss: 0.3631 - val_accuracy: 0.8445
Epoch 4/50
2910/2910 [=====] - 7s 2ms/step - loss: 0.3580 - accuracy: 0.8463 - val_loss: 0.3606 - val_accuracy: 0.8460
Epoch 5/50
2910/2910 [=====] - 9s 3ms/step - loss: 0.3556 - accuracy: 0.8476 - val_loss: 0.3591 - val_accuracy: 0.8461
Epoch 6/50
2910/2910 [=====] - 7s 2ms/step - loss: 0.3541 - accuracy: 0.8475 - val_loss: 0.3581 - val_accuracy: 0.8467
Epoch 7/50
2910/2910 [=====] - 7s 2ms/step - loss: 0.3530 - accuracy: 0.8475 - val_loss: 0.3581 - val_accuracy: 0.8467
```

```
In [37]: history_df2 = pd.DataFrame(history2.history)

plt.plot(history_df2.loc[:, ['accuracy']], "#BDE2E2", label='Training accuracy')
plt.plot(history_df2.loc[:, ['val_accuracy']], "#C2C4E2", label='Validation accuracy')

plt.title('Training and Validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



Overfitting and underfitting are checked using training and validation accuracy, as well as other metrics.

As the accuracy of training and validation varied slightly, our model is not overfitting. The fact that our model's training and validation accuracy are both quite good shows that it is learning patterns in the data.

```
In [39]: model2.evaluate(X_test, y_test)

910/910 [=====] - 3s 3ms/step - loss: 0.3487 - accuracy: 0.8467
```

```
Out[39]: [0.348657488822937, 0.8466588854789734]
```

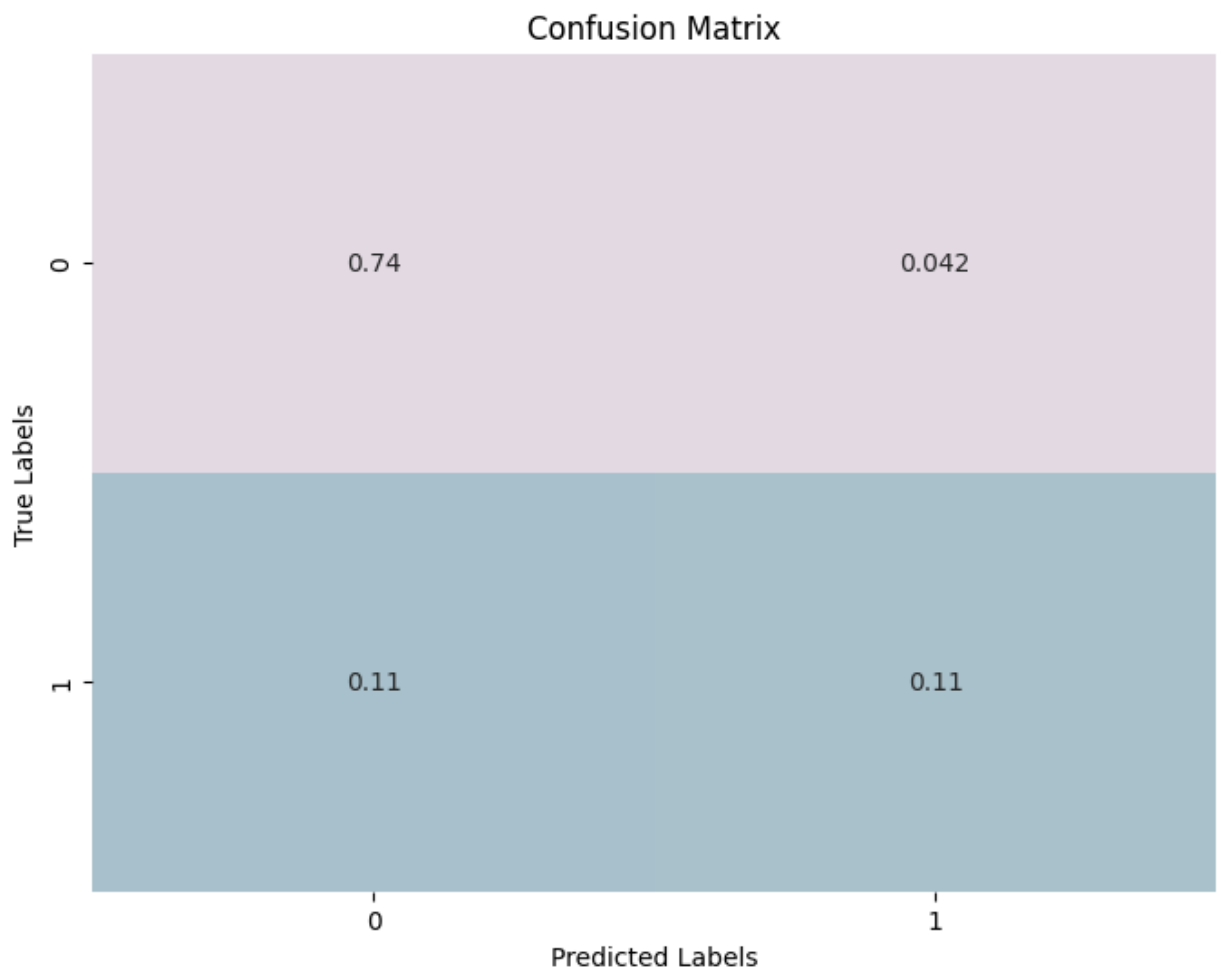


```
In [40]: y_predict = model2.predict(X_test)
y_pred_classes = (y_predict > 0.5)
```

```
910/910 [=====] - 1s 1ms/step
```

```
In [41]:
```

```
cmap = 'twilight'
# Plot the confusion matrix using Seaborn
cm = confusion_matrix(y_test, y_pred_classes)
plt.figure(figsize=(8, 6))
sns.heatmap(cm/np.sum(cm), annot=True, cmap=cmap, cbar=False)
plt.xlabel("Predicted Labels")
plt.ylabel("True Labels")
plt.title("Confusion Matrix")
plt.show()
```



```
In [42]: report = classification_report(y_test, y_pred_classes)
print(report)
```

	precision	recall	f1-score	support
0	0.87	0.95	0.91	22672
1	0.72	0.50	0.59	6420
accuracy			0.85	29092
macro avg	0.80	0.72	0.75	29092
weighted avg	0.84	0.85	0.84	29092

```
In [ ]:
```

A confusion matrix is a tabular representation that summarizes the performance of a classification model. It provides insights into how well the model's predictions match the actual class labels. The matrix displays true positive, true negative, false positive, and false negative counts, enabling assessment of accuracy, precision, recall, and F1-score. It's a valuable tool for understanding a model's strengths and weaknesses in predicting different classes.

```
In [ ]:
```

Future Works

Data Quality Enhancement: Focus on data cleaning, validation, and augmentation to ensure the quality of input data, which directly impacts the accuracy of predictions.

Performance Evaluation: Conduct thorough validation and comparison with existing rainfall prediction models to assess the effectiveness and superiority of our approach.

Longer Time Horizons: Extend the prediction horizon to forecast rainfall on longer time scales, such as monthly or seasonal predictions, to cater to different planning and decision-making needs.

Advanced Machine Learning Techniques: Explore more advanced techniques, or hybrid model to futhure improve prediction accuracy.

```
In [ ]:
```

Thank you

In []: